

Computer Architecture (Practical Class)

Assembly: Stack and Functions - Parameters and Local Variables

Luís Nogueira

Departamento de Engenharia Informática
Instituto Superior de Engenharia do Porto

lmn@isep.ipp.pt

2021/2022

- Many (if not most) functions require some sort of input data
- In x86-64, most of the data passed to procedures take place via registers
- Up to six integral (i.e., integer and pointer) arguments can be passed via registers
- When a function has more than six integral arguments, the other ones are passed on the stack
- This method works equally as well for any Assembly program, even if it wasn't derived from a C program

- The registers are used in a **specified order**, with the name used for a register depending on the size of the data type being passed

Operand size (bits)	Argument number					
	1	2	3	4	5	6
64	%rdi	%rsi	%rdx	%rcx	%r8	%r9
32	%edi	%esi	%edx	%ecx	%r8d	%r9d
16	%di	%si	%dx	%cx	%r8w	%r9w
8	%dil	%sil	%dl	%cl	%r8b	%r9b

- Consider the following functions:

Function f1

```
long f1(long a, long b, long c){  
    return a + b + f2(c);  
}
```

Function f2

```
long f2(long c){  
    return c + 1;  
}
```

- How can you implement an equivalent code in Assembly?

Passing parameters via registers: Example (2/2)

Function f1

```
long f1(long a, long b, long c){  
    return a + b + f2(c);  
}
```

Function f2

```
long f2(long c){  
    return c + 1;  
}
```

Function f1

```
f1:  
# a in %rdi, b in %rsi, c in %rdx  
# a + b  
addq %rsi, %rdi  
# store %rdi on stack  
pushq %rdi  
# prepare function call f2(c)  
movq %rdx, %rdi  
call f2  
# restore %rdi  
popq %rdi  
# a + b + f2(c)  
addq %rdi, %rax  
ret
```

Function f2

```
f2:  
# c in %rdi  
# c + 1  
incq %rdi  
movq %rdi, %rax  
ret
```

- Consider the following functions:

Function f1

```
void f1(int a, char b, int *res){  
    *res = f2(a,b);  
}
```

Function f2

```
int f2(int a, char b){  
    return a * b;  
}
```

- Implement an equivalent code in Assembly

Practice problem

Function f1

```
void f1(int a, char b, int *res){
    *res = f2(a,b);
}
```

Function f2

```
int f2(int a, char b){
    return a * b;
}
```

Function f1

```
f1:
    # a in %edi, b in %sil
    # *res in %rdx
    # store %rdx on stack
    pushq %rdx
    # f2(a,b)
    call f2
    #restore %rdx
    popq %rdx
    movl %eax, (%rdx)
    ret
```

Function f2

```
f2:
    # a in %edi, b in %sil
    # sign extend b
    movsbl %sil, %eax


    # a * b
    # result in %edx:%eax
    imull %edi
    ret
```

Passing parameters via stack (1/2)

- Assume that procedure P calls procedure Q with n integral arguments, such that $n > 6$
- P , the caller, must copy arguments 1 to 6 into the appropriate registers, and put arguments 7 through n onto the stack, with **argument 7 at the top of the stack**
- When passing parameters on the stack, all data sizes are rounded up to be multiples of eight
- With the arguments in place, procedure P can then execute a *call* instruction to transfer control to procedure Q
 - When the *call* instruction is executed, it places the return address from P onto the top of the stack as well, so Q knows where to return
- Q , the callee, can access its first six arguments via registers and the remaining ones from the stack

Passing parameters via stack (2/2)


- All these details leave the registers and the stack in the state shown here



Accessing parameters on the stack(1/2)

- The stack pointer (RSP) points to the top of the stack, where the return address is located
 - All of the input parameters for the function are located “underneath” the return address on the stack
- Popping values off of the stack to retrieve the input parameters would cause a problem, as the return address might be lost in the process
- Each parameter can be indirectly accessed via the offset from the RSP register, without having to POP values off of the stack
 - `movq 8(%rsp), %rax` - copies (without popping!) the 7th parameter to the RAX register

- Using indirect addressing with the RSP register to access the 7+ input parameters




Using RBP as the base pointer (1/2)

- There is an inconvenient problem with this approach! What happens if data is pushed to the stack throughout the callee's code?
- If this happens, it would change the location of the RSP stack pointer and throw off the indirect addressing values for accessing the parameters in the stack
- To avoid this problem, it is common practice to copy the RSP register value to the RBP register when entering the function
- This ensures that there is a register that always contains the correct pointer to the top of the stack when the function is called (a push does not change RBP)
- To avoid corrupting the RBP register on the caller, the callee (before the RSP register value is copied) should also place the RBP register value on the stack (RBP is a callee-saved register) and restore it before returning

Using RBP as the base pointer (2/2)

- All these details leave the stack in the state shown here



- This approach has created a standard set of instructions that are found in procedures that make a heavy use of the stack
- Stack space is allocated in **frames**
 - A section of stack used by one procedure call to store context while running, delimited by *%rbp* and *%rsp*
- Procedure code manages stack frames explicitly
 - **Prologue/Setup**: allocate space at start of procedure
 - **Epilogue/Cleanup**: deallocate space before return

- The base pointer *%rbp* being a stable “anchor” to the beginning of the stack frame throughout the execution of a function, is very convenient for manual Assembly coding and for debugging

Function template with Prologue and Epilogue


```
function:
    # prologue
    pushq %rbp          # save the original value of RBP
    movq  %rsp, %rbp    # copy the current stack pointer to RBP

    # function body
    ...

    #epilogue
    movq  %rbp, %rsp    # retrieve the original RSP value
    popq  %rbp          # restore the original RBP value
    ret
```

Prologue and Epilogue (3/3)

- The prologue saves the caller's value for *%rbp* into the callee's stack frame and the current stack pointer in *%rbp*
- This adjusted value of *%rbp* is the callee's frame pointer
 - The callee will not change this value until just before it returns
 - The frame pointer provides a stable reference point for local variables and caller arguments
- The epilogue ensures that the caller's *%rsp* and *%rbp* are restored, which deallocates all callee's reserved temporary space



Cleaning out the stack (1/2)

- Before the function (with more than 6 arguments) is called, the caller places the 7+ required input values in the stack
- When the called function returns, those values are still on the stack (since the callee function accessed them without popping them off of the stack)
- The caller should remove the old input values from the stack to get the stack back to where it was before the function call
 - This is particularly important if other values are going to be popped out of the stack after the callee's return
 - Ensures that the *ret* instruction at the end of the caller finds the proper return value on top of the stack
- While you can use the POP instruction to do this, you can also just move the RSP stack pointer back to the original location before the function call

Cleaning out the stack (2/2)

- Adding back the size of the data elements pushed onto the stack using the ADD instruction clears all the pushed data with a single instruction

Cleaning the Stack after Return

```
function:
...

pushq %rax      # push 8th parameter onto stack
pushq %rbx      # push 7th parameter onto stack
call utilfunc
addq $16, %rsp  # get the stack back to where it was before the function call

...
ret
```

Passing parameters via stack: Example (1/3)

- Consider the following functions:

Function f1

```
long f1(long a, long b, long c,  
        long d, long e, long f){  
    return a + f2(a*b,b,c,d,e,f,10,-20);  
}
```

Function f2

```
long f2(long a, long b, long c,  
        long d, long e, long f,  
        long g, long h){  
    /* do someting with parameters */  
    ...  
    return h/g;  
}
```

- How can you implement an equivalent code in Assembly?

Passing Parameters

```
#long f1(a,b,c,d,e,f)
f1:
    # a in %rdi, b in %rsi, c in %rdx
    # d in %rcx, e in %r8, f in %r9

    pushq %rdi                # save %rdi on stack

    movq %rdi, %rax           # a * b
    imulq %rsi

                                # prepare function call
    movq %rax, %rdi           # a*b in %rdi
    pushq $-20                # 8th parameter
    pushq $10                 # 7th parameter
    call f2                   # f2(a*b,b,c,d,e,f,10,-20);
    addq $16, %rsp            # clear 7th and 8th parameters

    popq %rdi                 # restore %rdi
    addq %rdi, %rax            # a + f2(...)
    ret
```

Passing parameters via stack: Example (3/3)

Accessing Parameters

```
#long f2(a,b,c,d,e,f,g,h)
f2:
    pushq %rbp                #prologue
    movq %rsp, %rbp

    # a in %rdi, b in %rsi, c in %rdx
    # d in %rcx, e in %r8, f in %r9
    # g in 16(%rbp), h in 24(%rbp)

    pushq %rbx                # save %rbx

    # do something with parameters
    ...
    movq 16(%rbp), %rbx        # get g from stack
    movq 24(%rbp), %rax        # get h from stack
    cqto                      # h/g
    idivq %rbx

    popq %rbx                 # restore %rbx

    movq %rbp, %rsp          # epilogue
    popq %rbp
    ret
```

- At times, local data must be stored in memory. Common cases of this include:
 - There are not enough registers to hold all of the local data
 - The address operator '&' is applied to a local variable, and hence we must be able to generate an address for it
 - Some of the local variables are arrays or structures and hence must be accessed by array or structure references
- When looking for an easy location for working storage for data elements within a function, the stack again comes to the rescue

Reserving space for local variables (1/3)


- The idea is to allocate storage for local variables when entering the function and release it before returning
- This space should be reserved after setting the RBP register to point to the top of the stack without affecting how the function's parameters are accessed
- Typically, a function allocates space on the stack frame by decrementing the stack pointer (%rsp)

Reserving Space for Local Variable

```
function:
    # prologue
    pushq %rbp
    movq  %rsp, %rbp
    subq  $16, %rsp    # reserve 16 bytes for local variables
                       # this space can be used, for e.g., as either
                       # four 32-bit values or two 64-bit values
    ...
```

Reserving space for local variables (2/3)

- By keeping `%rbp` pointing to the initial position throughout the execution of the function, we can refer local variables and parameters with fixed offsets relative to `%rbp`



- Now, if any data is pushed onto the stack, it will be placed after the local variables
- Local variables are then preserved and they can still be accessed via RBP
- When the epilogue of the function is reached and the RSP register is set back to its original value, the local variables will be lost from the stack
- They will not be directly accessible using the RSP or RBP registers from the calling function (thus the term “local variables”)

Using local variables: Example (1/3)

- Consider the following functions:

Function swap_add

```
long swap_add(long *xp, long *yp){  
    long x = *xp;  
    long y = *yp;  
    *xp = y;  
    *yp = x;  
    return x + y;  
}
```

Function f1

```
long f1(){  
    long arg1 = 100;  
    long arg2 = 200;  
    long sum = swap_add(&arg1,&arg2);  
    long diff = arg1 - arg2;  
    return sum * diff;  
}
```

- How can you implement an equivalent code in Assembly?

Using local variables: Example (2/3)

Function swap_add

```
#long swap_add(long *xp, long *yp)
swap_add:
    movq (%rdi), %rax    # get x
    movq (%rsi), %rcx    # get y
    movq %rcx, (%rdi)    # *xp = y
    movq %rax, (%rsi)    # *yp = x
    addq %rcx, %rax      # x + y
    ret
```

Using local variables: Example (3/3)

Function f1

```
#long f1()
f1:
    pushq %rbp                # prologue
    movq %rsp, %rbp
    subq $16, %rsp            # allocate 16 bytes for local variables

    movq $100, -8(%rbp)        # arg1 = 100;
    movq $200, -16(%rbp)       # arg2 = 200;

                                # prepare function call
    leaq -8(%rbp), %rdi        # compute &arg1 as first argument
    leaq -16(%rbp), %rsi       # compute &arg2 as second argument
    call swap_add              # call swap_add(&arg1,&arg2)

    movq -8(%rbp), %rdx        # get arg1
    subq -16(%rbp), %rdx       # diff = arg1 - arg2
    imulq %rdx                 # sum * diff

    movq %rbp, %rsp           # epilogue (deallocates 16 bytes)
    popq %rbp
    ret
```