

Computer Architecture (Practical Class)

Heterogeneous Data Structures

Luís Nogueira

Departamento de Engenharia Informática
Instituto Superior de Engenharia do Porto

lmn@isep.ipp.pt

2022/2023

In C, we have two ways of combining variables of different types:

- *structures*, declared using the keyword `struct`, aggregate multiple variables under the same name;
- *unions*, declared using the keyword `union`, allow a single variable to use several data types.

Today, we will discuss *structures in C and Assembly*.

- A *structure* is:
 - A continuous memory region (similarly to arrays), composed of different *members*
 - Members are accessed by their name and can have different types
- We can have structures in C by:
 - defining a single structure in a variable;
 - declaring a structure type;
 - declaring a new data type.

In the example below, we define a *single structure* in a variable called “c1”.

Structure in C - Defining a structure held by a variable

```
#include <stdio.h>

/* define a structure in a global variable c1 */
struct {
    int    n_wheels;
    int    n_passengers;
    float  fuel_consumption;
} c1;

int main(){
    /* init and access the data in c1 */
    c1.n_wheels=4;
    c1.n_passengers=2;
    c1.fuel_consumption=12.5;
    printf("The car has %d wheels, %d passengers.Consumption is %f l/100 km.\n",
        c1.n_wheels, c1.n_passengers, c1.fuel_consumption);
    return 0;
}
```

- The structure has only one instantiation (in c1); no other instantiations of the structure exist.

Defining structures in C (2/3)

In the example below, we declare a *structure type* called “s_car”.

Structure in C - Declaring a structure type

```
#include <stdio.h>

/* declare a structure type called s_car */
struct s_car {
    int    n_wheels;
    int    n_passengers;
    float  fuel_consumption;
} ;

int main(){
    /* define two structures of type s_car */
    struct s_car c1, c2;

    /* init and access the data in c1 and c2 */
    c1.n_wheels=4; c1.n_passengers=2; c1.fuel_consumption=12.5;
    c2.n_wheels=5; c2.n_passengers=4; c2.fuel_consumption=20.0;
    printf("Car 2 has %d wheels, %d passengers. Consumption is %f l/100 km.\n",
        c2.n_wheels, c2.n_passengers, c2.fuel_consumption);
    return 0;
}
```

- To instantiate the *s_car structure*, we define new variables using the name of the structure after the keyword `struct`.

Defining structures in C (3/3)

In the example below, we declare a new *data type* called “t_car” using the keyword “typedef”.

Structure in C - Declaring a data type for a structure

```
#include <stdio.h>

/* declare a data type called t_car */
typedef struct { /* we could write 'typedef struct s_car' */
    int    n_wheels;
    int    n_passengers;
    float  fuel_consumption;
} t_car; /* this is the name of the structure data type */

int main(){
    /* define two cars */
    t_car c1, c2;

    /* init and access the data in c1 and c2 */
    c1.n_wheels=4; c1.n_passengers=2; c1.fuel_consumption=12.5;
    c2.n_wheels=5; c2.n_passengers=4; c2.fuel_consumption=20.0;
    printf("Car 2 has %d wheels, %d passengers. Consumption is %f l/100 km.\n",
           c2.n_wheels, c2.n_passengers, c2.fuel_consumption);
    return 0;
}
```

- To instantiate the *t_car* data type, we define new variables using the name of the data type.

Pointers to structures in C (1/2)

- We can define pointers to structures

```
struct s_car *ptr;    /* pointer to structure type */  
t_car *ptr;           /* pointer to a structure data type */
```

- To access the fields of a structure referenced by a pointer, use ->

```
ptr->n_wheels  
ptr->n_passengers  
ptr->fuel_consumption
```

instead of:

```
(*ptr).n_wheels  
(*ptr).n_passengers  
(*ptr).fuel_consumption
```

Notes:

- To obtain the address of a structure, use the & operator
- Structures are very similar to arrays (continuous block of memory), **except in this aspect** (must use & to get address)

Structure in C - Obtaining a pointer to a structure

```
#include <stdio.h>

typedef struct {
    int    n_wheels;
    int    n_passengers;
    float  fuel_consumption;
} t_car;

int main(){
    /* define a car */
    t_car c1;
    /* define a pointer to a car */
    t_car *car_ptr = &c1;

    /* init and access the data in c1 using the pointer */
    car_ptr->n_wheels=4;
    car_ptr->n_passengers=2;
    car_ptr->fuel_consumption=12.5;
    printf("The car has %d wheels, %d passengers. Consum. is %f l/100 km.\n",
           car_ptr->n_wheels, car_ptr->n_passengers, car_ptr->fuel_consumption);

    return 0;
}
```

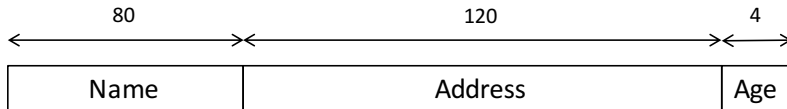
Storing structures in memory

- Consider the following structure in C

Example structure in C to store personal data

```
struct s_person_data{ /* we could also have used typedef ... */  
    char name[80];      /* 80 bytes */  
    char address[120];  /* 120 bytes */  
    int age;            /* 4 bytes */  
};
```

- To store this structure, we need 204 bytes ($80+120+4$)



Important note

We will see some subtleties of structure sizes in a moment

- In Assembly, we refer to members within a structure by their *offset*
- The offset of each member of the structure refers to the displacement, in bytes, from the start of the structure
- To simplify this access, we can define constants for the offset of each member using the `.equ` directive

```
.equ DATA_SIZE, 204      # total size
.equ NAME_OFFSET, 0       # name is at the beginning of the structure
.equ ADDRESS_OFFSET, 80   # the address starts at byte 80
.equ AGE_OFFSET, 200      # age starts at byte 200
```

Function set_age in C

```
void set_age(struct  
    s_person_data *pdata){  
    /* age = 30 */  
    pdata->age = 30;  
}
```

Function set_age in Assembly

```
set_age:  
    # *pdata on %rdi  
    # age = 30  
    movl $30, AGE_OFFSET(%rdi)  
    ret
```

Example: Access to a structure member

Structure Declaration	Access in C	Access in Assembly
<pre>struct rec{ int x; int a[3]; int *p; };</pre>	<pre>void set_x(struct rec *r, int val){ /* access by name */ r->x=val; }</pre>	<pre># *r in %rdi # val in %esi # access by offset movl %esi, (%rdi)</pre>

Memory layout for the structure rec

x	a	p
0 ... 3	4 ... 15	16 ... 23

- Note that the address of *x* (the first member of the structure) is equal to the address of *r* (the structure)

Example: Obtaining the address of a structure member

Structure Declaration	Access in C	Access in Assembly
<pre>struct rec{ int x; int a[3]; int *p; }</pre>	<pre>int* find_a(struct rec *r, unsigned int i){ /* address of a[i] */ return &(r->a[i]); }</pre>	<pre># *r in %rdi # i in %esi # %rax = r + 4 + (4 * i) leaq 4(%rdi,%rsi,4),%rax</pre>

Memory layout for the structure rec

x	a	p
0 ... 3	4 ... 15	16 ... 23

- To access field *a*, we need to add the appropriate offset to the address of the structure (in this case 4 bytes). Then, we need to access the element *i* within *a*, by adding offset $i * 4$

Example: Obtaining a pointer and changing a structure member

Structure Declaration	Access in C	Access in Assembly
<pre>struct rec{ int x; int a[3]; int *p; }</pre>	<pre>void set_p(struct rec *r){ /* address of a[i], where i is the value of member x */ r->p = &(r->a[r->x]); }</pre>	<pre># *r in %rdi # %rcx = r->x movslq (%rdi),%rcx # %rax = r + 4 + 4*(r->x) leaq 4(%rdi,%rcx,4),%rax # r->p = %rax movq %rax,16(%rdi)</pre>

Memory layout for the structure rec

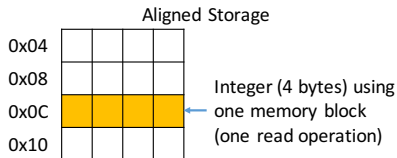
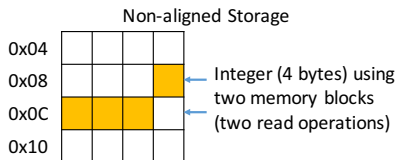
x	a	p
0 ... 3	4 ... 15	16 ... 23

- Based on the size of each member (according to its data type), and the start address of the structure, we must compute the addresses

Why

- Memory is accessed in blocks of fixed size
 - Usually, 4 or 8 bytes, depending on the system
- Storing data at addresses multiple of the block size — *aligned to the block size* — will reduce the number of memory accesses

The following example depicts a system accessing memory in blocks of 4 bytes:



Data alignment is performed by the compiler

- **The compiler inserts “spaces” in the data stored in memory to ensure that the members are aligned**
- This improves the performance of the code
- This is required in some architectures; in x86-64 it is just recommended
 - in x86-64 data can be misaligned, at the cost of a performance penalty
- May be treated differently by different operating systems

The alignment rules are based on the principle that any primitive object of K bytes must have an address that is a multiple of K

- $K = 1$ byte: **char**
 - No restrictions
- $K = 2$ bytes: **short**
 - The least significant bit must be 0 (that is, the address must be a multiple of 2)
- $K = 4$ bytes: **int, float, ...**
 - The 2 least significant bits must be 0 (that is, the address must be a multiple of 4)
- $K = 8$ bytes: **double, long, long long, char *, ...**
 - The 3 least significant bits must be 0 (that is, the address must be a multiple of 8)

Important notes

These restrictions hold for most x86-64 operating systems, except that on Windows, the *long* type has size and alignment 4 (The *long long* type has size and alignment 8 on all x86-64 operating systems)

Particularly important to know about data alignment when dealing with structures in C

- In order to be able to share structures between C and Assembly, we have to be aware of the data alignment made by the compiler

Inside the structure

- We must satisfy the alignment requirements for each member of the structure
- Each structure has an alignment requirement of \mathbb{K} , which may require implicit internal padding, depending on the previous member

Placement of the structure in memory

- Given \mathbb{K} , the largest alignment requirement inside the structure:
- The starting address of the structure must be a multiple of \mathbb{K}
- The total size of the structure must be a multiple of \mathbb{K} , which may require padding after the last member (external padding)

Data alignment in structures: Examples

Total size: 24 bytes

- $\mathbb{K}=8$, due to the member *c* of type long

```
struct S1{  
    char a;  
    int b[2];  
    long c;  
};
```

a	3 bytes	b	4 bytes	c
0	[gap]	4 ... 11	[gap]	16 ... 23

Total size: 24 bytes

- $\mathbb{K}=8$, due to the member *b* of type long

```
struct S2{  
    int a;  
    long b;  
    short c;  
};
```

a	4 bytes	b	c	6 bytes
0 ... 3	[gap]	8 ... 15	16 .. 17	[gap]

Important note

The starting address of both structures must be a multiple of $\mathbb{K} = 8$

Consider the following data type:

```
typedef struct{
    char age;
    short number;
    int grades[10];
    char name[80];
    char address[120];
}Student;
```

- Develop in Assembly the function `void save_grades(Student *s, int *new_grades, int size)` that copies all the elements of the array `new_grades` to the field `grades` of the structure pointed by `s`

Function save_grades in Assembly

```
# void save_grades(Student *s, int *new_grades, int size)
save_grades:
    # *s in %rdi, *new_grades in %rsi, size in %edx

    addq $4, %rdi                # rdi = &(s->grades[0])
    movslq %edx, %rdx            # rdx = size
    cmpq $0, %rdx
    jle end

    movl $0, %rcx                # index = 0

loop_grades:
    cmpq %rcx, %rdx
    je end

    movl (%rsi, %rcx, 4), %eax    # place grade to copy on %eax
    movl %eax, (%rdi, %rcx, 4)    # copy grade in %eax to new_grades

    incq %rcx                    # index++
    jmp loop_grades

end:
    ret
```